FINDING SIMILAR ITEMS

LinkedIn Jobs & Skills

Magdalena Skowerska

Declaration

I declare that this material, which I now submit for assessment, is entirely my own work and has not been taken from the work of others, save and to the extent that such work has been cited and acknowledged within the text of my work, and including any code produced using generative AI systems. I understand that plagiarism, collusion, and copying are grave and serious offences in the university and accept the penalties that would be imposed should I engage in plagiarism, collusion or copying. This assignment, or any part of it, has not been previously submitted by me or any other person for assessment on this or any other course of study.

Table of Contents

Declaration *1*

Abstract *4*

Dataset *5*

Finding Similar Items Techniques *6*

Data preprocessing and Tokenization 6

TF- IDF 7

Cosine distance measure 9

Locality Sensitive Hashing (LSH) 10

K-Nearest Neighbors (k-NN) 11

Conclusions *12*

FIGURE 1 - NUMBER OF JOB DESCRIPTIONS DUPLICATES 7
FIGURE 2 - IDF IMPLEMENTATION 8
FIGURE 3 - OPTIMIZED IDF IMPLEMENTATION 8
FIGURE 4 - CUSTOM COSINE SIMILARITY IMPLEMENTATION RESULTS 10
FIGURE 5 - SKLEARN COSINE SIMILARITY IMPLEMENTATION RESULTS 10
FIGURE 6 - MINHASHLSH RESULTS FOR SMALL DATASET 10
FIGURE 7 - MINHASHLSH RESULTS FOR LARGE DATASETS 10
FIGURE 8 - K-NN RESULTS 11

Abstract

The aim of this project is to find similar pairs or sets of job offers published on Linkedin. The column used to build the similarity detector is job_summary in the job_summary.csv file.
Different techniques will be considered and analized from the performance perspective and the dataset size. TF-IDF technique will be chosen, implemented and analized with the cosine distance measure by using custom functions, sklearn library and pyspark library. Locality Sensitive Hashing (LSH) will be implemented by using datasketch library for small datasets and it will be used in combination with the TF-IDF algorithm for big datasets. Finally, for demonstrative purposes, there will be proposed an implementation of K-Nearest Neighbors (k-NN) suitable for small dataset.

Dataset

The chosen datset is the ‘ Linkedin Jobs & Skills’ for the year 2024, scraped from LinkedIn and published on Kaggle under the ODC-By license. It is composed by three files:

- job_skills.csv containing links to the job offer and the related skills.
- job_summary.csv containing links to the job offer and its description.
- linkedin_job_posting.csv containing job details, company information, locations etc.

The files are in the csv format with tab-separated-values and formatted in the UTF-8 character set.

In this project we will use also the dataset ‘stopwords’ from the python library Natural Language Toolkit (NLTK) which contains stopwords in the English language such as: ‘the’, ‘are’, ‘is’ etc. , and the dataset ‘wordnet’ which groups words into sets of synonyms.

Finding Similar Items Techniques

In order to find similar documents or items in natural language processing (NLP), there are different techniques that can be considered such as:

1. Count vectorization: consisting in transforming a collection of text documents into a matrix where each row corresponds to a document and each columns to te unique token values. The value of the matrix cells represent the count of a specific word in a specific document.
2. Term Frequency times Inverse Document Frequency (TF-IDF): similar to count vectorization method but it adjusts the weights of the words on the basis of their frequency across all the documents.
3. Word Embeddings like Word2Vec, GloVe (Global Vectors for Word Representation) or FastText: Represent words as dense vectors in a continuous vector space, capturing semantic relationships.
4. Locality Sensitive Hashing (LSH): based on the usage of the hash function to create buckets of documents that result similar with high probability.

The documents we are dealing with are job descriptions, which can be considered short and of known context, advanced techniques like Word Embeddings will be not considered since they are relatively computationally expensive. Instead, the TF-IDF technique will be implemented that, contrary to the Count vectorization, takes into account the importance of words in relation to a given document making it a more sophisticated approach.

Data preprocessing and Tokenization

The data considered for applying the different techniques comes from the job_summary table consisting of two columns: job_link and job_summary.

As a first step, we apply the data preprocessing techniques including:

- Checking and removing missing values: for both columns there are not null values.
- Checking the precence of duplicates: for the column 'job_link' the number of distinct values is equal to the number of total rows. However, for the column 'job_summary' the number of distinct value is lower than the number of total rows. In fact, as we can see below the column contains duplicates that will be kept in the dataset since they are referring to different job links.

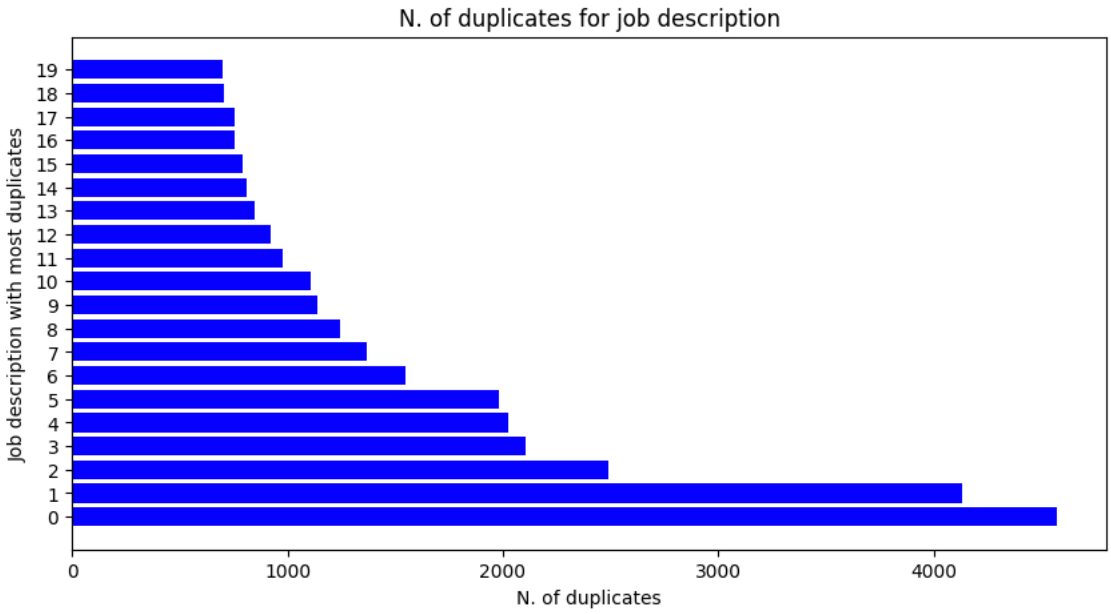

Figure 1 - number of job descriptions duplicates

The second step consists of the tokenization process creating for each job offer a list of words contained in its description. To improve performance and limit the noise in the data we will remove the stopwords from the lists and convert all tokens to lowercase. Furthermore, we will use also the stemming and lemmatizing process considering nltk library:

- Stemming is the process of removing suffixes and prefixes from a word to reduce it to its root form.
- Lemmatization is the process of reducing a word to its base form by considering the sets of synonyms provided by 'wordnet' dataset.

TF- IDF

Term Frequency times Inverse Document Frequency (TF-IDF) consists of calculating:

- **The Term Frequency Index**: given 't' the occurences of a token in a document and 'd' the total number of tokens in a document, the Term Frenquency of the considered token is:

$$TF(t,d) = \frac{t}{d}$$

- **The Inverse Document Frequency**: given a token 't', 'U' the corpus of documents, 'N' the number of documents contained in 'U' and n(t) the number of documents containing the token considered, The Inverse Document Frequency is:

$$IDF(t) = \frac{N}{n(t)}$$

- **TF-IDF** measure of the token in the document considered is:

$$TF(t,d) * IDF(t)$$

First, the algorithm will be implemented by using a custom function. Below there are shown two implementations of the function that is used to compute the Inverse Document Frequency values for all of the tokens in the corpus.

```python
def idfs(corpus):
    """ Compute IDF
      Args:
          corpus (RDD): input RDD of documents represented by job link and description
      Returns:
          dictionary: dictionary of tokens to IDF values
    """

    unique_tokens = corpus.flatMap(lambda doc: doc[1]).distinct().collect()
    N = corpus.count()
    idfs_val = {token: 0 for token in unique_tokens}

    for token in unique_tokens:
        doc_freq = corpus.filter(lambda doc: token in doc[1]).count()
        if doc_freq > 0:
          idfs_val[token] = N / doc_freq
        else:
          0

    return idfs_val
```

Figure 2 - idf implementation

```python
def idfs_optimized(corpus):
    """ Compute IDF
      Args:
          corpus (RDD): input RDD of documents represented by job link and description
      Returns:
          RDD: RDD of tokens to IDF values
    """

    unique_tokens = corpus.flatMap(lambda doc: doc[1]).distinct().collect()
    N = corpus.count()

    token_doc_pairs = corpus.flatMap(lambda doc: [(token, doc[0]) for token in set(doc[1])])
    doc_freq_optimized = token_doc_pairs.map(lambda pair: (pair[0], 1)).reduceByKey(lambda a, b: a + b)
    idfs_val = doc_freq_optimized.map(lambda token_freq_record: (token_freq_record[0], N / token_freq_record[1] if token_freq_record[1] > 0 else 0))

    return idfs_val
```

Figure 3 - optimized idf implementation

When considering a sample dataset of 10 job links, the improvement of the idfs to idfs_optimized function will reduce the computation time from 10 minutes to only 4 seconds. This is done by replacing the filter() and count() functions with flatMap() and reduceByKey(). In fact, the first approach:

- Requires two passes through the data. First, the filter function iterates over all the document to check the presence of the token, then the count function iterates again through the data to count the filtered records.
- Can lead to higher memory usage because the intermediate results of the filtering process must be stored until the count operation is complete.

The second approach on the other hand:

- Generates pairs in a single pass and then aggregates them using MapReduce approach.

The optimized version will be used in the tfidf_optimized function that allows all computation to be done exploiting the RDD data structure suitable for huge amount of

data. On the other hand the function tfidf is based on the use of dictionaries. So, the limitation of the non optimized function is the number of tokens. Since the tfidf function returns a dictionary of all tokens and their tf-idf values considering all documents, if the number of tokens does not fit in the main memory the application of the algorithm will fail. In our case the documents are short (considering the first 10000 records the maximum length description is 2234) so we can assume that all tokens of one document can fit into the main memory.

It is possible to use the TF-IDF algortihm with a library scikit-learn built on top of Numpy, SciPy, and matplotlib. This is done by using TfidfVectorizer which is a class that converts a collection of documents into a matrix with one row per document, all the unique tokens as columns and the TF-IDF scores as its values. This class, however, allows in input only list of strings which can be acceptable in case of small datasets but not with larger ones since it brings all data in memory of the driver leading to memory issues. Furthermore, it produces in output a sparse matrix which can result not efficient from a storage perspective.

For very large dataset we need to consider the use of the distributed computing system. In this case we can use our custom optimized function or the machine learning library provided by pyspark. In particular, we will use the CountVectorizer that given all the documents produces a list of unique tokens, counts the number of times the tokens appear in the document, and produces a matrix with one document per row, the list of tokens and a vector representing the indices of non-zero entries with the relative counts. The result of the computation will be given to the IDF model that calculates the IDF values for each token and produce the TF-IDF vectors for each document. This method is efficient from the memory usage perspective since it stores the sparse matrix as sparse vectors made of: the total number of non-zero elements, their values and their indices.

Cosine distance measure

The cosine similarity is based on measuring the cosine of the angle between two vectors in a multi-dimensional space. A smaller angle (closer to 1) means the documents are more similar.
We define our own function that calculate the cosine similarity between two job descriptions given in input an RDD data structure containing job links with the relative dictionary document tokens to its TF-IDF values and a list of the job links for which we want to find similar job descriptions. This function will be used for the optimized verion of the custom TF-IDF algorithm. The memory limitation of this function can be represented by the list of the target job links and the length of the tokenized description which makes it suitable for large datasets.
We consider also the use the sklearn implementation when building TF-IDF matrix with the function TfidfVectorizer().fit_transform() that can be used in case of small datasets.

Below there is the sample results by applying the custom implementation and sklearn.

```
The 2 most similar link jobs to https://www.linkedin.com/jobs/view/restaurant-manager-at-rock-n-roll-sushi-3805551344 are:
  - https://au.linkedin.com/jobs/view/experienced-sushi-chef-at-an-authentic-japanese-restaurant-at-world-mode-holdings-3788943496
  - https://www.linkedin.com/jobs/view/store-manager-in-training-at-texas-thrift-3775401861

The 2 most similar link jobs to https://www.linkedin.com/jobs/view/med-surg-registered-nurse-rn-at-touchette-regional-hospital-3732389852 are:
  - https://www.linkedin.com/jobs/view/registered-nurse-flex-at-magnolia-regional-health-center-3734113140
  - https://www.linkedin.com/jobs/view/rn-4-tower-at-magnolia-regional-health-center-3648967253

The 2 most similar link jobs to https://www.linkedin.com/jobs/view/registered-nurse-cath-lab-at-stonesprings-hospital-center-3799543261 are:
  - https://www.linkedin.com/jobs/view/rn-cath-lab-sign-on-bonus-at-grand-strand-medical-center-%E2%80%93-hca-3801579946
  - https://www.linkedin.com/jobs/view/registered-nurse-cardiac-cath-lab-at-texas-nursing-services-3787762317
```

Figure 4 - custom cosine similarity implementation results

```
The most similar link jobs to 'https://www.linkedin.com/jobs/view/restaurant-team-leader-3582-highway-114-ft-worth-tx-unit-%231119-at-whataburger-3671804517'
    based on the threshold 0.5 are:
 - https://www.linkedin.com/jobs/view/restaurant-team-leader-2859-n-germantown-pkwy-germantown-tn-unit-%231300-at-whataburger-3650208124
```

Figure 5 - sklearn cosine similarity implementation results

Locality Sensitive Hashing (LSH)

LSH is used for efficient candidate generation. It hashes similar items (like document vectors) into the same bucket with high probability so that the number of comparisons is reduced to those elements that are in the same bucket.
The similarity measure is done by using Jaccard similarity that compares two sets of tokens contained in the documents by dividing the size of their intersection to the size of their union.
As for the other algorithms the project presents the implementation for a small dataset of 10 records and the implementation necessary in case of huge amount of data using pyspark library. In the second case, the algorithm can be performed on the TF-IDF vectors obtained by applying the pyspark IDF model. To retrieve the most similar items given a link job the approxNearestNeighbors() method will be used based on the cosine similarity between the vectors hashed to the same bucket. The number of hash tables is 5 and for each hash table the number of buckets is set to the default value equal to 128, the number of hash functions (by the parameter numHashFunctions). The MinHashLSH, as TF-IDF one, from pyspark operates over spark DataFrame built on top of RDD data structure with some optimizations thanks to the Catalyst optimizer and Tungsten engine.

Below the sample results obtained by applying the algorithm.

```
The most similar link jobs to 'https://www.linkedin.com/jobs/view/restaurant-team-leader-3582-highway-114-ft-worth-tx-unit-%231119-at-whataburger-3671804517'
      based on a threshold 0.5 are:
 - https://www.linkedin.com/jobs/view/restaurant-team-leader-2859-n-germantown-pkwy-germantown-tn-unit-%231300-at-whataburger-3650208124

The most similar link jobs to 'https://www.linkedin.com/jobs/view/restaurant-team-leader-2859-n-germantown-pkwy-germantown-tn-unit-%231300-at-whataburger-3650208124'
      based on a threshold 0.5 are:
 - https://www.linkedin.com/jobs/view/restaurant-team-leader-3582-highway-114-ft-worth-tx-unit-%231119-at-whataburger-3671804517
```

Figure 6 - MinHashLSH results for small dataset

```
The 2 most similar link jobs to 'https://www.linkedin.com/jobs/view/restaurant-manager-at-rock-n-roll-sushi-3805551344' are:
 - https://www.linkedin.com/jobs/view/kitchen-manager-at-talent-up-recruiting-llc-3771487939
 - https://www.linkedin.com/jobs/view/kitchen-manager-at-electric-3795348717

The 2 most similar link jobs to 'https://www.linkedin.com/jobs/view/med-surg-registered-nurse-rn-at-touchette-regional-hospital-3732389852' are:
 - https://www.linkedin.com/jobs/view/surgical-technologist-at-ssm-health-3707557565
 - https://www.linkedin.com/jobs/view/electroneurodiagnostic-technologist-at-ssm-health-3741459082

The 2 most similar link jobs to 'https://www.linkedin.com/jobs/view/registered-nurse-cath-lab-at-stonesprings-hospital-center-3799543261' are:
 - https://www.linkedin.com/jobs/view/rn-cath-lab-sign-on-bonus-at-grand-strand-medical-center-%E2%80%93-hca-3801579946
 - https://www.linkedin.com/jobs/view/texas-roadhouse-kitchen-manager-at-texas-roadhouse-3742056109
```

Figure 7 - MinHashLSH results for large datasets

K-Nearest Neighbors (k-NN)

K-nn is a supervised machine learning algorithm based on the fact that similar data points are likely to be close to each other in the feature space and they are grouped in k groups. It uses the cosine distance to measure the distance between points and it identifies the k, in this case equal to two, nearest points (neighbors).

Below we can see the results of the algorithm for our dataset of 10 records:

```
The most similar link job to 'https://www.linkedin.com/jobs/view/restaurant-manager-at-rock-n-roll-sushi-3805551344' is:
- https://www.linkedin.com/jobs/view/restaurant-team-leader-2859-n-germantown-pkwy-germantown-tn-unit-%231300-at-whataburger-3650208124 (Distance: 0.8413)

The most similar link job to 'https://www.linkedin.com/jobs/view/med-surg-registered-nurse-rn-at-touchette-regional-hospital-3732389852' is:
- https://www.linkedin.com/jobs/view/registered-nurse-cath-lab-at-stonesprings-hospital-center-3799543261 (Distance: 0.7846)

The most similar link job to 'https://www.linkedin.com/jobs/view/registered-nurse-cath-lab-at-stonesprings-hospital-center-3799543261' is:
- https://www.linkedin.com/jobs/view/med-surg-registered-nurse-rn-at-touchette-regional-hospital-3732389852 (Distance: 0.7846)

The most similar link job to 'https://uk.linkedin.com/jobs/view/commercial-account-executive-at-the-recruit-lab-3805254225' is:
- https://uk.linkedin.com/jobs/view/concession-store-manager-selfridges-at-linda-farrow-3796128638 (Distance: 0.8478)

The most similar link job to 'https://www.linkedin.com/jobs/view/store-manager-at-stop-shop-3782135496' is:
- https://www.linkedin.com/jobs/view/restaurant-team-leader-3582-highway-114-ft-worth-tx-unit-%231119-at-whataburger-3671804517 (Distance: 0.8582)

The most similar link job to 'https://www.linkedin.com/jobs/view/restaurant-team-leader-3582-highway-114-ft-worth-tx-unit-%231119-at-whataburger-3671804517' is:
- https://www.linkedin.com/jobs/view/restaurant-team-leader-2859-n-germantown-pkwy-germantown-tn-unit-%231300-at-whataburger-3650208124 (Distance: 0.0222)

The most similar link job to 'https://www.linkedin.com/jobs/view/hair-stylist-meadow-brook-center-at-jobs-for-humanity-3786387856' is:
- https://uk.linkedin.com/jobs/view/commercial-account-executive-at-the-recruit-lab-3805254225 (Distance: 0.8743)

The most similar link job to 'https://uk.linkedin.com/jobs/view/concession-store-manager-selfridges-at-linda-farrow-3796128638' is:
- https://uk.linkedin.com/jobs/view/commercial-account-executive-at-the-recruit-lab-3805254225 (Distance: 0.8478)

The most similar link job to 'https://www.linkedin.com/jobs/view/material-handler-at-intelliswift-software-3789502416' is:
- https://www.linkedin.com/jobs/view/restaurant-team-leader-3582-highway-114-ft-worth-tx-unit-%231119-at-whataburger-3671804517 (Distance: 0.9162)

The most similar link job to 'https://www.linkedin.com/jobs/view/restaurant-team-leader-2859-n-germantown-pkwy-germantown-tn-unit-%231300-at-whataburger-3650208124' is:
- https://www.linkedin.com/jobs/view/restaurant-team-leader-3582-highway-114-ft-worth-tx-unit-%231119-at-whataburger-3671804517 (Distance: 0.0222)
```

Figure 8 - k-nn results

Conclusions

This project considers different techniques that could be used to implement an algorithm for finding similar items between set of documents. In our case, the documents are represented by the job descriptions and the TF-IDF algorithm was chosed. More than one way of implmementation were proposed. Among them the ones suitable for big datasets with a large amount of records are: the optimized custom version and the implementation using pyspark library. For demonstrative purposes it was also shown the possibility to use K-nn and MinHashLSH with small datasets.